

Practical SQL

Session 01

Tuesday 01 October 2013

General Assembly

Welcome

- Class focus is SQL syntax
- We will also cover:
 - Enough background to orient thinking
 - Usage of SQL to explore and analyze data
- Exercises away from class build familiarity

Who are you and why are you here?

- What sort of companies do you work for?
- What are your roles?
- What are you using for analysis?
- What are you interested in?
 - Replicating Excel analyses with SQL
 - Running queries
 - Orientation to SQL and database concepts

Who Am I

- Entrepreneur / consultant
 - Data management
 - CRM migration and integration
- Self-taught on LAMP(ython) stack
 - Plus some DNS, bash scripting, XML / XSLT
- Prior professional experience:
 - Office of Management and Budget
 - Investment banking (JP Morgan, UBS, boutique)
 - MBA, University of Chicago

Roadmap – Session 1

- Preliminaries
- What is SQL
 - What it is used for
 - What is meant by it
- Software and sample database installation
- Loading data from CSV
- Query syntax basics
 - Running queries
 - Filtering, grouping
 - Functions
 - Combining tables
- Using queries for analysis

Tools

- Database Software
 - Some means of running SQL queries is needed
 - MySQL, Postgres, Sql Server, Oracle
 - Microsoft Access is NOT recommended
 - If you have questions about your SQL implementation, see me
 - MySQL is free, download from mysql.com
 - MySQL
 - Both “Community Server” and startup package
 - MySQL Workbench
- Text editor – you may or may not need one
 - Wordpad, Text Edit are very basic examples
 - Textmate, Sublime Text are more featured tools
 - Windows: Notepad++
 - Do NOT use a word processor (e.g. Word)
 - Do NOT use ‘rich text format’ / *.rtf
 - SQL code should just be ASCII text, with NO formatting information

Get A Database Program

- Some means of running SQL queries is needed
 - MySQL, Postgres, Sql Server, Oracle
 - Microsoft Access is NOT recommended
 - If you have questions about your SQL implementation, see me
 - MySQL is free, download from mysql.com
 - MySQL
 - Both “Community Server” and startup package
 - MySQL Workbench

What is SQL -- Usage

Role of SQL in Data Analysis

- Data analysis requires:
 - Getting, cleaning, loading, ‘modeling’
 - Exploring, manipulating, extracting, analyzing
 - Reporting, visualization
- SQL is best at analysis, but call help with all of this

Getting and Cleaning

- Something else has to get the data to CSV or database
- SQL string functions can help clean and reformat data
- Tricks with LOAD and INSERT can remodel data

Reporting and Visualization

- SQL has zero visualization capacity, but it can output data as needed for your graphics program
- “Reporting” is here sort of a term of art
 - E.g. pivot tables, cross tabs, and so forth
 - SQL can do an awful lot, but isn’t great at this
 - Implementations do provides base layers to reporting tools like Jasper, Pentaho, etc

Analysis and Transformation

- SQL's stock in trade is filtering, aggregating and grouping on set-based criteria
 - Filter WHERE records meet some stated criteria
 - State equals Illinois, name in (Smith, Baker, Han)
 - Date before 2012-06-01
 - Score greater than 15
 - GROUP BY values in a field, or in a combination of fields
- Uses that capacity to combine information in different sets / tables by matching records on some criteria
- We can nest these abilities into “subqueries”, generating criteria on the fly

As A Language – Application Usage

- I.e., why we care
- A syntax for:
 - Describing tables
 - Extracting and manipulating data from those tables
 - Plus dealing with the craft of managing data: loading, indexing, etc
- The action is in bullet #2, but we have to deal with the other stuff too

Extracting and Manipulating

- Extracting:
 - We want to specify criteria for which records we display or analyze
 - Combine data extracted from several tables
- Manipulating:
 - Transform values in fields
 - Calculate values for fields of groups of records
- These fundamental operations can be extended to very complex transformations

What is SQL -- Architecture

A “Standard”, Not a Program

- There are many SQLs
 - MySQL / Postgres / SQLite / Microsoft Access / Oracle / Sybase / etc etc etc
- Each is an implementation of a standard
 - “SQL” is basically a program’s promise to behave in a defined way in response to a particular query for a particular set of data

“Client / Server”

- Most implementations are actually two programs
 - One is a front end / client, giving us access to a server and showing its outputs
 - MySQL Workbench
 - The other is a back end / server, holding the data and processing the queries
 - MySQL
 - MySQL Workbench is to MySQL as Chrome browser is to an HTTP server

Client / Server Implications

- We can access existing data
 - We don't need data dumps to do analysis
 - We do need some sensible and safe access to the existing databases
 - If your production database has a million records, you can analyze them all without downloading them all
- We can access remote data
 - If your servers are on Rackspace or Amazon you can still get at the data
- Programs can access the data
 - Languages like Ruby, Python, PHP all have tools for interacting with databases
 - They act as clients, but getting data for (say) serving web pages rather than display to an analyst

Excel vs SQL

Why Excel?

- Good graphic interface orients us to the data and the tools
 - We can find tools just by clicking around the menus
- Excel organizes models visually, so we can see the relationships among cells
 - We can see the cells flowing into a formula
 - We imagine extensions of our models by thinking through new visual relationships
- Excel handles many issues we don't even know about
 - Until we use a tool with less hand-holding – like SQL

Why SQL? -- Analytic

- Operated through text statements
 - Easy to identify and extend patterns
 - Replicable
 - If a new query resembles a prior one, just copy and revise the prior
 - Chainable
 - We can easily add conditions and “filters”
 - We can write a query and refer to its results in another query
 - Traceable
 - We can see how we got to a result by examining the query
- Non-visual modeling allows us to fluidly work in more than two dimensions

Implications

- Excel gets used like a Swiss Army Knife
 - We have problems, Excel gives us path to solve them
- SQL can do the same things and more, BUT requires infrastructure, training and problem solving – i.e., brain damage
 - When we do need it, who has time for training?

Why SQL? -- Organizational

- A lot of data already in your company's relational databases
- Dev team familiar with SQL
- SQL easier to automate
- Easy to operate over a network

Getting Started

Getting Files

- Download files from https://github.com/chernevik/practical_sql
 - Click 'Download ZIP' to get an archive of the files
 - Open / unpack archive

Install Database Software

- If you don't have access to a database, follow instructions in files
- Note that you must install both a client program and a server program
 - Probably, also, a script for starting the server on computer boot
- Reach out to me if you have trouble

Install Sample Database

- In files, find script 'setup_sampdb.sql'
- Just run this in your MySQL client
 - This is MySQL specific; tell me what you are using and I'll provide scripts to load from CSV
- If you read this file, it's just a big query that creates a database and tables, and inserts data

Queries

Run a Query

- In your client, type:
 - ‘SELECT 2 + 2;’ and ‘execute’
- Note that queries are synonymous with ‘commands’
 - We use them to get information about the databases and tables available
 - We can test out some functions and expressions with them

Queries to Explore Syntax

- SQL evaluates more sophisticated expressions, like logic:
 - `SELECT 1 AND 2 -> 1`
 - `SELECT 2 AND 2 -> 1`
 - `SELECT 0 AND 2 -> 0`
- It supports functions:
 - `SELECT LEFT('cat in the hat', 3) -> 'cat'`
 - `SELECT RIGHT('cat in the hat', 3) -> 'hat'`
 - `SELECT UPPER('cat') -> 'CAT'`
 - `SELECT LENGTH('cat in the hat') -> 14`
 - `SELECT IF(1, 'OK', 'NOT') -> 'OK'`
 - `SELECT IF(0, 'OK', 'NOT') -> 'NOT'`
- Each of these character strings is a SQL 'statement', made up of 'expressions'
- The database program evaluates each expression, and 'returns' a result for it
 - Those results may become input to another, enclosing expression
 - `SELECT IF(0 and 2, 'OK', 'NOT') -> 'NOT'`
- We rarely use SQL this way, but it does show how expressions and functions work

How do we retrieve data?

- So:
 - ‘SELECT * FROM sampdb.student’
 - This query / statement gets all fields from the table student in the database sampdb
 - We will add more clauses to this to control its output

Queries are text, period

- Whether entered by typing, or loaded from a file, or loaded in some other way, the client only cares that it has text to pass to the server
- For example, in the downloaded materials, find:
 - sessions/01/simple.sql
 - Open in query editor
 - Execute

Aside: Organizing Your Environment

- Queries are typically written in text files
 - Again, no word processing, no .doc or .rtf files
- Make a place on your computer to hold your files
 - data, queries, drafts
 - Your files can expand quickly, as you copy files to experiment or to make variations
 - Use of subdirectories helps keep things organized
 - Code repository tools like Git and Mercurial make a lot of sense
 - This would be a great time to begin learning and using these

Loading Data

Annoying but necessary

- No fundamental concepts involved here – just file management
- Understanding this process does orient us to the fundamentals of what the computer is doing
- We will go over this fast and revisit in exercises
 - Next week we'll go over problems in those

First, get data to database

- Get data into CSV format
 - Excel can do this for us
 - On Mac:
 - “Comma Separated Values” will have CR line terminators
 - “Windows CSV” will have CRLF terminators
- Line terminators? Hang on

Create a database

- Simple query creating a home for our table

```
DROP DATABASE IF EXISTS learning;  
CREATE DATABASE learning;
```

Create a table

- Very simple table
- Notice
homogeneity of
data type (all
varchar)

```
DROP TABLE IF EXISTS learning.from_excel;  
CREATE TABLE learning.from_excel (
```

```
    date varchar(20),  
    category varchar(20),  
    event_id varchar(20)  
);
```

Load data

- Let's break this down

- CR =
‘carriage
return’
- LF = ‘line
feed’

LOAD DATA LOCAL

INFILE "/Users/gimli/Desktop/from_excel_windows.csv"
INTO TABLE learning.from_excel

FIELDS

TERMINATED BY ','
ENCLOSED BY ''

LINES

TERMINATED BY '\r\n' -- for files with CRLF terminators
-- TERMINATED BY '\r' -- for CR terminators

IGNORE 1 LINES;

A moment on file data structure

- Loading CSV to SQL forces us to consider how our data is represented in a file
- Excel handles a variety of file storage formats for us – it will guess which is the proper one
- SQL will not
 - It is a power user tool
 - Stuff like file format has defaults but if these don't work we must figure things out for ourselves

File data structure: Problem

- The file is a string of bytes – in our simple case, ASCII data
 - Unicode? Wrong class
- Some of those bytes are data, some are metadata
 - How does MySQL know where one field begins and another ends?
 - Where one record ends another begins?

Solved by convention and definition

- Fields “delimited”, or “terminated”, by commas
 - When MySQL sees a comma, it starts a new field
- Fields are also “enclosed” by double quotes
 - When MySQL sees a double quote, it considers all that follows a field until it sees another double quote
- Not redundant
 - This structure allows us to have commas in our data
 - If an enclosure is missing MySQL will just go with the field terminator
 - Double quotes in data can be handled by escaping mechanism (e.g. \" is not an enclosure) but we haven't time for that

Solved by convention and definition - 2

- Lines (e.g. records, rows) terminated by a line termination character(s)
 - In Windows, generally carriage return + line feed, or CR LF ('\\r\\n')
 - In Mac and Linux, line feed ('\\n')
- These characters are VERY ANNOYING because we almost never see them
 - Our programs all have facilities for interpreting them for us
- We can figure this out or just guess

NEVER MIND -- JUST DO THIS

- Create a database
- Create a table
- Save as CSV
- Run this
- Check results, if bad fiddle with this and try again

```
LOAD DATA LOCAL
  INFILE "/Users/gimli/Desktop/from_excel_windows.csv"
  INTO TABLE learning.from_excel

FIELDS
  TERMINATED BY ','
  ENCLOSED BY '"'

LINES
  TERMINATED BY '\r\n' -- for files with CRLF terminators
  -- TERMINATED BY '\r' -- for CR terminators

IGNORE 1 LINES;
```

Query Syntax

Simple Query

```
SELECT name, sex, student_id  
FROM sampdb.student;
```


- Gets all fields of all records in table

Let's unpack that:

- “database”
 - A program concept for holding “tables”
 - Comparable to a folder in your computer's file system
- “table”
 - A program concept holding data about instances of a particular entity type – students, presidents, exams
 - These data are held in “records” or “rows”
- “record”
 - Data about a particular instance – about a particular student or particular movie
 - Data is in a collection of fields or columns
 - Comparable to a row in a spreadsheet
- “Column” / “field”
 - An attribute of an entity type about which data is collected – a student's name, or sex, or the state of origin of a president
 - Comparable to a cell in a spreadsheet row

Simple Query

```
SELECT name, sex, student_id  
FROM sampdb.student;
```



We could also use:

```
SELECT * FROM sampdb.student
```

- << * >> is a wildcard meaning "all columns"

Filtering the results

```
SELECT  
  *  
FROM  
  sampdb.student  
WHERE  
  sex = "F";
```

- Notice the reorganization of the text for clarity
- WHERE clause defines conditions imposed on returned rows

An Aside on “query files”

- MySQL clients are looking for text only – e.g. files containing only ASCII or Unicode
- Don't use Word, don't use “rich text format”
 - These introduce characters that will confuse the MySQL client
- Use TextEdit, SublimeText, MacVim
 - vi !!!
 - A “text editor” of some sort

Filtering the results by 'tuples'

```
SELECT
    first_name,
    last_name,
    state

FROM
    sampdb.president

WHERE
    (last_name, state) IN

    (
        ('Adams', 'MA')
    )
;
```

- A 'tuple' is a set of fields
 - A table is actually a set of tuples
- The IN operator allows us to see if a particular group of fields has a particular set of values

Analyzing the results

```
SELECT  
    COUNT(student_id)  
FROM  
    sampdb.student  
;
```

- COUNT() is a “function”
- Functions return information about or manipulate contents of a field
- Some functions apply only to field of single record
- Others, like COUNT, apply to the field for multiple records
- COUNT returns the number of not NULL values in the field
 - We’re ignoring “NULL” for now

Analyzing the results -- Aggregation

```
SELECT
    sex,
    COUNT(student_id),
    COUNT(sex)
FROM
    sampdb.student
GROUP BY
    sex
;
```

- COUNT() is an a “aggregate” function
- We can segment the records on which it works
- GROUP BY clause gathers up records by their value in a given field
- The function can be applied to any field, not just that by which records are grouped

Analyzing the results -- Prettify

```
SELECT
    sex AS Gender,
    COUNT(student_id) AS "ID Count",
    COUNT(sex) AS gender_count
FROM
    sampdb.student
GROUP BY
    gender
    -- sex

    -- this is a comment
    -- all on this line after "--" is ignored

    # so is this, for MySQL
    /*
        And this, for MySQL
    */
;
```

- Inside the column expressions we've added "as [name]" clauses
- These are "alias" terms
- The results aren't changed but their display is
- We can use the alias to specify the column

Functions

```
SELECT
    first_name,
    last_name,
    state,
    ROUND(DATEDIFF(death, birth) / 365, 1)
AS age

FROM
    sampdb.president

ORDER BY
    age DESC

;
```

- DATEDIFF() is a defined MySQL function returning the number of days between two dates
- ROUND() controls decimal points
 - (we need this here to make some later code work)

More Functions, Aggregated

```
SELECT
  state,
  MAX(
    ROUND(DATEDIFF(death, birth) / 365, 1)
  )
AS max_age,
  COUNT(*)

FROM
  sampdb.president

GROUP BY
  state

ORDER BY
  max_age DESC

;
```

- MAX() is an aggregate function, returning the largest of a set of values
- Here the set is the values calculated for a group of rows
- GROUP BY defines the rows
- We drop the name fields because they aren't meaningful when rows are GROUPed

Find Me

- gordon@practicalhorseshoeing.com
- Twitter: @gordonagress
- github.com/chernevik/practical_sql
- 917-620-3402