

Practical SQL

Operations

We can do simple arithmetic math

```
1 SELECT
2     score / 100 as pct_score,
3     CAST(score AS decimal(8, 2)) / 100 as pct_score_2
4 FROM
5     sampdb.score
6 LIMIT 5;
7
```

Clear

Run

```
1 +-----+-----+
2 | pct_score | pct_score_2 |
3 +-----+-----+
4 |    0.2000 |    0.200000 |
5 |    0.2000 |    0.200000 |
6 |    0.1800 |    0.180000 |
7 |    0.1300 |    0.130000 |
8 |    0.1800 |    0.180000 |
9 +-----+-----+
10
```

Clear

Run

-- Note that MySQL will automatically convert the value returned as a decimal.

-- Other SQL versions will not. Postgres requires a type conversion of the INT to a DECIMAL type.

More complex math

We can do more complex math

```
1 SELECT
2     student_id,
3     score,
4     ( score - 15.1 ) / 3.8 AS z_score,
5     ( CAST(score AS DECIMAL(8, 2) ) - 15.1 ) / 3.8 AS z_score
6     -- ( point - mean ) / std. dev
7 FROM
8     sampdb.score
9 WHERE
10     event_id = 1
11 LIMIT 5
12 ;
13
```

[Clear](#)[Run](#)

```
1 +-----+-----+-----+-----+
2 | student_id | score | z_score | z_score |
3 +-----+-----+-----+-----+
4 |          1 |    20 |  1.28947 |  1.289474 |
5 |          3 |    20 |  1.28947 |  1.289474 |
6 |          4 |    18 |  0.76316 |  0.763158 |
7 |          5 |    13 | -0.55263 | -0.552632 |
8 |          6 |    18 |  0.76316 |  0.763158 |
9 +-----+-----+-----+-----+
10
```

[Clear](#)[Run](#)

We are 'hard coding' the mean and standard deviation values

We'll discuss how to get these on the fly

Strings

String functions manipulate the appearance of text

```
1 SELECT
2     first_name,
3     UPPER(first_name),
4     LOWER(first_name)
5
6 FROM
7     sampdb.president
8
9 LIMIT 5;
10
```

Clear

Run

```
1 +-----+-----+-----+
2 | first_name | UPPER(first_name) | LOWER(first_name) |
3 +-----+-----+-----+
4 | George    | GEORGE            | george            |
5 | John      | JOHN              | john              |
6 | Thomas    | THOMAS            | thomas            |
7 | James     | JAMES             | james             |
8 | James     | JAMES             | james             |
9 +-----+-----+-----+
10
```

Clear

Run

String extraction

String functions can "reach into" strings

```
1 SELECT
2     last_name,
3     LENGTH(last_name) AS len,
4     LEFT(last_name, 3) AS lft,
5     RIGHT(last_name, 3) AS rght,
6     SUBSTRING(last_name, 3, 2) AS sub,
7     SUBSTRING(last_name, LENGTH(last_name) - 2, 2) AS penultimate_2
8 FROM
9     sampdb.president
10 LIMIT 5;
11
```

[Clear](#)[Run](#)

```
1 +-----+-----+-----+-----+-----+-----+
2 | last_name | len | lft | rght | sub | penultimate_2 |
3 +-----+-----+-----+-----+-----+-----+
4 | Washington | 10 | Was | ton | sh | to |
5 | Adams      | 5  | Ada | ams | am | am |
6 | Jefferson   | 9  | Jef | son | ff | so |
7 | Madison     | 7  | Mad | son | di | so |
8 | Monroe      | 6  | Mon | roe | nr | ro |
9 +-----+-----+-----+-----+-----+-----+
10
```

[Clear](#)[Run](#)

Finding strings allows more interesting usages

String manipulation becomes more interesting when we start looking for strings

1
2
3
4
5
6
7
8

SELECT
 email,
 -- LOCATE('@', email),
 position('@' in email)
FROM
 sampdb.member
LIMIT 5;

ClearRun

1
2
3
4
5
6
7
8
9
10

+-----+-----+
| email | position('@' in email) |
+-----+-----+
jeanne_s@earth.com	9
august.lundsten@pluto.com	16
NULL	NULL
arbogast.ruth@mars.net	14
c.dorfman@uranus.net	10
+-----+-----+

ClearRun

Here we get the position of the @ sign

We can use that, with our other functions, to get a list of domains

Putting string functions together

With the location of the '@', we can extract the domain name

```
1 SELECT
2     email,
3     RIGHT(email, LENGTH(email) - LOCATE('@', email)) AS domain
4 FROM
5     sampdb.member
6 LIMIT 5;
7
```

Clear

Run

```
1 +-----+-----+
2 | email                | domain          |
3 +-----+-----+
4 | jeanne_s@earth.com   | earth.com       |
5 | august.lundsten@pluto.com | pluto.com      |
6 | NULL                 | NULL            |
7 | arbogast.ruth@mars.net | mars.net        |
8 | c.dorfman@uranus.net  | uranus.net      |
9 +-----+-----+
10
```

Clear

Run

Now, for each record, we can extract the domain name

Might use this to count members by domain our list use, or check for correlation with other factors

Date Functions

```
1 SELECT
2     birth,
3     death,
4     YEAR(birth) AS yr,
5     MONTH(birth) AS mnt,
6     MONTHNAME(birth) AS mnt_name,
7     DATEDIFF(death, birth) AS days,
8     DATEDIFF(death, birth) / 365 AS years
9 FROM
10     sampdb.president
11 LIMIT 5;
12
```

[Clear](#)[Run](#)

```
1 +-----+-----+-----+-----+-----+-----+-----+
2 | birth      | death      | yr  | mnt | mnt_name | days | years |
3 +-----+-----+-----+-----+-----+-----+-----+
4 | 1732-02-22 | 1799-12-14 | 1732 | 2   | February | 24767 | 67.8548 |
5 | 1735-10-30 | 1826-07-04 | 1735 | 10  | October  | 33119 | 90.7370 |
6 | 1743-04-13 | 1826-07-04 | 1743 | 4   | April     | 30397 | 83.2795 |
7 | 1751-03-16 | 1836-06-28 | 1751 | 3   | March     | 31150 | 85.3425 |
8 | 1758-04-28 | 1831-07-04 | 1758 | 4   | April     | 26729 | 73.2301 |
9 +-----+-----+-----+-----+-----+-----+-----+
10
```

[Clear](#)[Run](#)

Putting it together

We can combine records, operations and functions to calculate and format output

```
1 SELECT
2     CONCAT(
3         first_name,
4         " ",
5         last_name,
6         " lived ",
7         ROUND(DATEDIFF(death, birth) / 365, 1),
8         " years"
9     ) AS sentence
10 FROM
11     sampdb.president
12 LIMIT 5;
13
```

Clear Run

```
1 +-----+
2 | sentence
3 +-----+
4 | George Washington lived 67.9 years
5 | John Adams lived 90.7 years
6 | Thomas Jefferson lived 83.3 years
7 | James Madison lived 85.3 years
8 | James Monroe lived 73.2 years
9 +-----+
10
```

Clear Run

More on Aggregate Functions

We attack this through dataset exploration

SELECT DISTINCT gives us the set of values in an attribute -- no repetitions

```
1 SELECT DISTINCT
2     state
3 FROM
4     sampdb.president
5 ORDER BY
6     state
7 ;
8
```

Clear Run

```
1 +-----+
2 | state |
3 +-----+
4 | AR    |
5 | CA    |
6 | CT    |
7 | GA    |
8 | IA    |
9 | IL    |
10 | KY    |
11 | MA    |
12 | MO    |
13 | NC    |
14 | NE    |
15 | NH    |
16 | NJ    |
17 | NY    |
18 | OH    |
19 | PA    |
20 | SC    |
21 | TX    |
22 | VA    |
23 | VT    |
24 +-----+
25
```

Clear Run

This is basically what we get from GROUP BY

GROUP BY subdivides the dataset into subsets

Each subset is characterized by having the same value for the GROUP BY attribute

```
1 SELECT
2     state,
3     COUNT(state) AS count
4 FROM
5     sampdb.president
6 GROUP BY
7     state
8 ORDER BY
9     state
10 ;
11
```

Clear

Run

```
1 +-----+-----+
2 | state | count |
3 +-----+-----+
4 | AR   | 1     |
5 | CA   | 1     |
6 | CT   | 1     |
7 | GA   | 1     |
8 | IA   | 1     |
9 | IL   | 1     |
10 | KY   | 1     |
11 | MA   | 4     |
12 | MO   | 1     |
13 | NC   | 2     |
14 | NE   | 1     |
15 | NH   | 1     |
16 | NJ   | 1     |
17 | NY   | 4     |
18 | OH   | 7     |
19 | PA   | 1     |
20 | SC   | 1     |
21 | TX   | 2     |
22 | VA   | 8     |
23 | VT   | 2     |
24 +-----+-----+
25
```

Clear

Run

Same analysis, against member

```
1 SELECT
2     state,
3     COUNT(state) AS count
4 FROM
5     sampdb.member
6 GROUP BY
7     state
8 ORDER BY
9     state
10 LIMIT 15;
11
```

[Clear](#)[Run](#)

```
1 +-----+-----+
2 | state | count |
3 +-----+-----+
4 | AK    | 1     |
5 | AL    | 3     |
6 | AZ    | 1     |
7 | CA    | 6     |
8 | CO    | 3     |
9 | CT    | 1     |
10 | FL    | 5     |
11 | GA    | 3     |
12 | HI    | 1     |
13 | IA    | 2     |
14 | ID    | 1     |
15 | IL    | 6     |
16 | IN    | 3     |
17 | KS    | 2     |
18 | KY    | 2     |
19 +-----+-----+
20
```

[Clear](#)[Run](#)

DISTINCT also works by _tuples_

```
1 SELECT DISTINCT
2   city,
3   state
4 FROM
5   sampdb.member
6 ORDER BY
7   state,
8   city
9 LIMIT 15
10 ;
11
```

Clear

Run

```
1 +-----+-----+
2 | city          | state |
3 +-----+-----+
4 | Fairbanks     | AK    |
5 | Dothan        | AL    |
6 | Huntsville    | AL    |
7 | Mobile        | AL    |
8 | Kayenta       | AZ    |
9 | Los Angeles   | CA    |
10 | Oakland       | CA    |
11 | San Francisco | CA    |
12 | Stockton      | CA    |
13 | Trona         | CA    |
14 | Denver        | CO    |
15 | Durango       | CO    |
16 | Waterbury     | CT    |
17 | Coral Gables  | FL    |
18 | Fort Myers    | FL    |
19 +-----+-----+
20
```

Clear

Run

Here we are getting a list of city / state value pairs that occur at least once in the data

This is not _permutations_, e.g. the possible combinations of these values -- each of these _does_ appear at least once

As does GROUP BY

Grouping by _several_ attributes subdivides counts

```
1 SELECT
2     city,
3     state,
4     COUNT(*)
5 FROM
6     sampdb.member
7 GROUP BY
8     state,
9     city
10 ORDER BY
11     state,
12     city
13 LIMIT 15
14 ;
15
```

Clear Run

```
1 +-----+-----+-----+
2 | city      | state | COUNT(*) |
3 +-----+-----+-----+
4 | Fairbanks | AK    | 1         |
5 | Dothan    | AL    | 1         |
6 | Huntsville| AL    | 1         |
7 | Mobile    | AL    | 1         |
8 | Kayenta   | AZ    | 1         |
9 | Los Angeles| CA   | 2         |
10 | Oakland   | CA    | 1         |
11 | San Francisco| CA  | 1         |
12 | Stockton  | CA    | 1         |
13 | Trona     | CA    | 1         |
14 | Denver    | CO    | 2         |
15 | Durango   | CO    | 1         |
16 | Waterbury | CT    | 1         |
17 | Coral Gables| FL   | 1         |
18 | Fort Myers| FL    | 1         |
19 +-----+-----+-----+
20
```

Clear Run

All the Alabama records are still in the output.

But the sum of Alabama records, which was 3, is now subdivided over three different cities

ORDER BY shows the grouping of the records

```
1 SELECT
2     last_name,
3     first_name,
4     member_id,
5     city,
6     state
7 FROM
8     sampdb.member
9 ORDER BY
10    state,
11    city
12 LIMIT 15
13 ;
14
```

[Clear](#)[Run](#)

	last_name	first_name	member_id	city	state
4	Matthews	Bill	56	Fairbanks	AK
5	Edwards	John	82	Dothan	AL
6	Hughes	Max	37	Huntsville	AL
7	Schauer	Alma	65	Mobile	AL
8	Kirby	Timothy	48	Kayenta	AZ
9	Puntillo	Cheryl	59	Los Angeles	CA
10	Pierson	Stanley	88	Los Angeles	CA
11	Sprague	Earl	99	Oakland	CA
12	Smith	Laura	98	San Francisco	CA
13	Feit	Daniel	19	Stockton	CA
14	Simmons	David	49	Trona	CA
15	Garner	Steve	89	Denver	CO
16	Sawyer	Dennis	7	Denver	CO
17	Bookstaff	Barbara	47	Durango	CO
18	Schenk	Cindy	42	Waterbury	CT

[Clear](#)[Run](#)

Referring to SQL results in SQL queries -- Subqueries

Subqueries allow us to embed one query within another

These embedded queries are 'subqueries'

We have to watch the table returned by the subquery carefully

The returned table has to have the fields and row expected of it by the calling query

Two varieties, 'correlated' and 'uncorrelated'

'correlated' is cooler but harder, we'll do 'uncorrelated' first

Presidents from the most presidential states

Say we want a list of those presidents from the three states that have sent the most presidents

We can get the list by a query:

```
1 SELECT
2     state,
3     COUNT(*) AS state_count
4 FROM
5     sampdb.president
6 GROUP BY
7     state
8 ORDER BY
9     state_count DESC
10 LIMIT 3
11 ;
12
```

Clear

Run

```
1 +-----+-----+
2 | state | state_count |
3 +-----+-----+
4 | VA    |           8 |
5 | OH    |           7 |
6 | MA    |           4 |
7 +-----+-----+
8
```

Clear

Run

Hard-code the values

We can put those values into a query

```
1 SELECT
2     last_name,
3     first_name,
4     state
5 FROM
6     sampdb.president
7 WHERE
8     state IN
9     (
10        'VA',
11        'OH',
12        'MA'
13    )
14 ORDER BY
15     state
16 ;
17
18
```

Clear Run

```
1 +-----+-----+-----+
2 | last_name | first_name | state |
3 +-----+-----+-----+
4 | Bush      | George H.W. | MA    |
5 | Adams     | John Quincy | MA    |
6 | Kennedy   | John F.     | MA    |
7 | Adams     | John       | MA    |
8 | Harding   | Warren G.   | OH    |
9 | Taft      | William H.  | OH    |
10 | McKinley  | William    | OH    |
11 | Harrison  | Benjamin    | OH    |
12 | Garfield  | James A.    | OH    |
13 | Hayes     | Rutherford B. | OH    |
14 | Grant     | Ulysses S.  | OH    |
15 | Taylor    | Zachary     | VA    |
16 | Tyler     | John        | VA    |
17 | Harrison  | William H.  | VA    |
18 | Monroe    | James       | VA    |
19 | Wilson    | Woodrow     | VA    |
20 | Madison   | James       | VA    |
21 | Jefferson | Thomas      | VA    |
22 | Washington | George     | VA    |
23 +-----+-----+-----+
24
```

Clear Run

But, Kludgy!

What if the data changes?

Joins

A basic query:

```
SELECT
    student_id AS Student,
    AVG(score) AS Average,
    COUNT(score) AS "# of Tests"
FROM
    sampdb.score
GROUP BY
    student_id
ORDER BY
    Average DESC
```

- Returns each student's average

	Student	Average	# of Tests	
▶	1	48.0000	5	
	27	45.7500	4	
	5	42.8333	6	
	18	41.3333	6	
	17	41.2000	5	
	2	40.4000	5	
	11	39.8333	6	

- But can we get the student's name rather than their id?

Yes

```
SELECT
  st.name AS      Name,
  scr.student_id AS Id,
  AVG(scr.score) AS Average,
  COUNT(scr.score) AS "# Tests"
FROM
  sampdb.score scr
INNER JOIN
  sampdb.student st
ON
  scr.student_id = st.student_id
GROUP BY
  scr.student_id
ORDER BY
  Average DESC
```

	Name	Id	Average	# Tests
▶	Megan	1	48.0000	5
	Carter	27	45.7500	4
	Abby	5	42.8333	6
	Max	18	41.3333	6
	Will	17	41.2000	5
	Josiah	2	40.4000	5

Breaking down the query

```
SELECT
  st.name AS      Name,
  scr.student_id AS Id,
  AVG(scr.score) AS Average,
  COUNT(scr.score) AS "# Tests"
FROM
  sampdb.score scr
INNER JOIN
  sampdb.student st
ON
  scr.student_id = st.student_id
GROUP BY
  scr.student_id
ORDER BY
  Average DESC
```

- A join combines rows from different tables
- The “ON” clause specifies conditions for combining records
 - Here, records with the same student_id value are combined into a single record
 - The SELECT clause still specifies which fields are displayed from that combined record
- INNER JOIN specifies which table to join data from
 - The addition of “scr” and “st” in the specification of tables provides alias values for reference by the other clauses
 - Without these the query wouldn’t know which table it should find a field in

Two Really Simple Tables

- T1

i1	c1
1	a
2	b
3	c

- t2

i2	c2
2	c
3	b
4	a

Inner Join: All Rows Matched

- A simple inner join matches each row in one table with each row in the other
- So joining a 3 row table with a 3 row table produces a 9 row table
- Obviously we don't want all those rows

```
SELECT * FROM join_sample.t1 INNER  
JOIN join_sample.t2;
```

i1	c1	i2	c2
1	a	2	c
2	b	2	c
3	c	2	c
1	a	3	b
2	b	3	b
3	c	3	b
1	a	4	a
2	b	4	a
3	c	4	a

Limiting Inner Join Results

```
SELECT
  *
FROM
  join_sample.t1
INNER JOIN
  join_sample.t2
WHERE
  join_sample.t1.i1 = join_sample.t2.i2
;
```

i1	c1	i2	c2
1	a	2	c
2	b	2	c
3	c	2	c
1	a	3	b
2	b	3	b
3	c	3	b
1	a	4	a
2	b	4	a
3	c	4	a

i1	c1	i2	c2
2	b	2	c
3	c	3	b

Left Join

- Get all rows from the “left” table, each with that row from the “right” table that matches on the specified fields

```
SELECT
  *
FROM
  join_sample.t1
LEFT JOIN
  join_sample.t2
ON
  join_sample.t1.i1 = join_sample.t2.i2
;
```

i1	c1	i2	c2
1	a	NULL	NULL
2	b	2	c
3	c	3	b

Contact

gordon@practicalhorseshoeing.com